

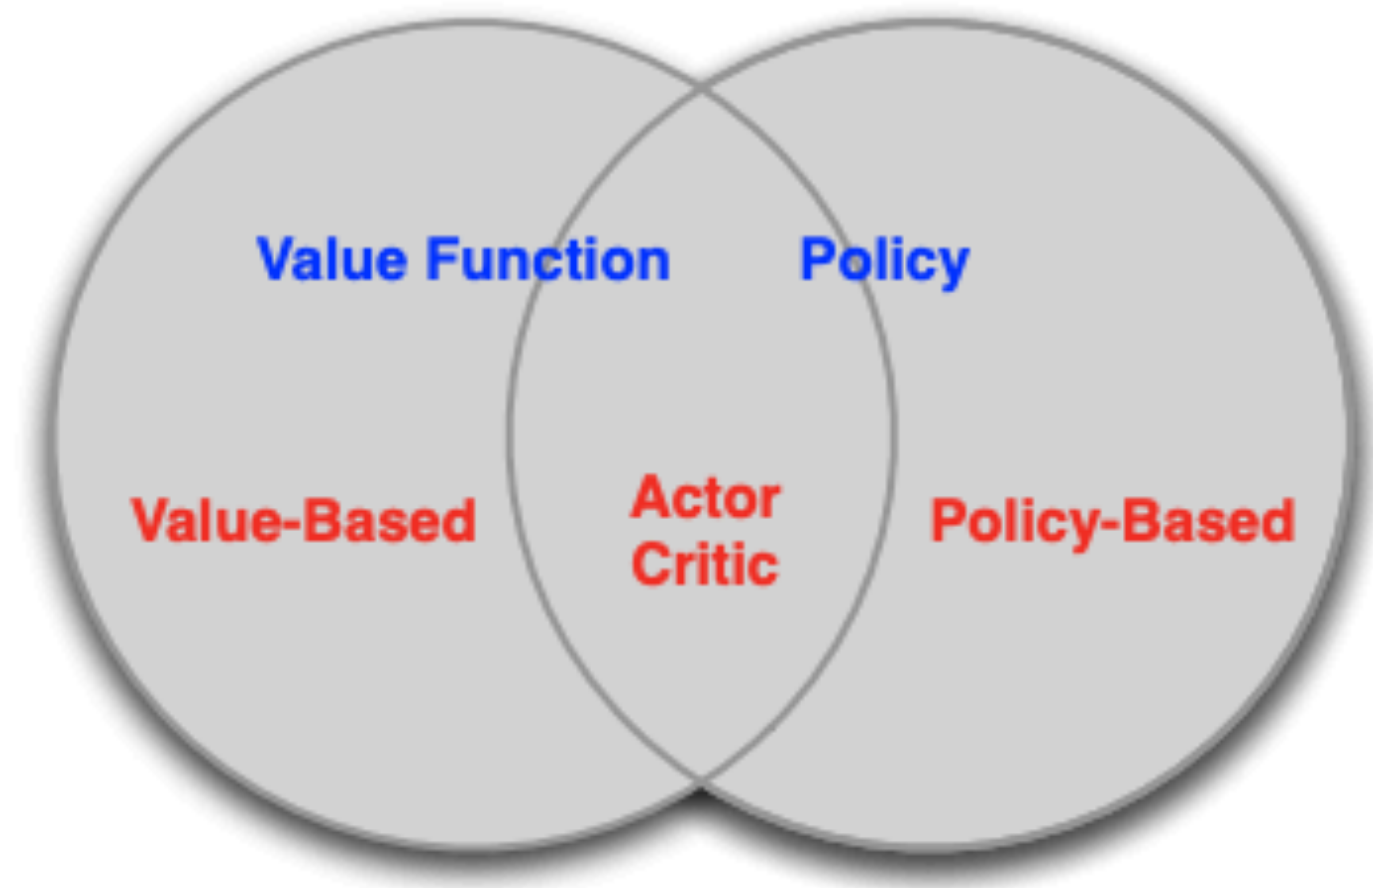
# **(Asynchronous) Actor Advantage Critics Reinforcement Learning: A2C and A3C**

Reinforcement Learning Group Meeting  
Yanjun Gao  
March 14, 2019




## Recall ... Value-Based and Policy-Based RL

- Value Based
  - Learnt Value Function
  - Implicit policy (e.g.  $\epsilon$ -greedy)
- Policy Based
  - No Value Function
  - Learnt Policy
- Actor-Critic
  - Learnt Value Function
  - Learnt Policy



Slides borrowed from: David Silver - Lecture 7: Policy Gradient Methods

## Recall ... Actor-Critic RL



# What is A2C and A3C? Why asynchronous?

## A2C: Actor Advantage Critics RL A3C: Asynchronous Actor Advantage Critics RL

*Classic actor-critics method with parallel architecture*

---

### Asynchronous Methods for Deep Reinforcement Learning

---

**Volodymyr Mnih**<sup>1</sup>

**Adrià Puigdomènech Badia**<sup>1</sup>

**Mehdi Mirza**<sup>1,2</sup>

**Alex Graves**<sup>1</sup>

**Tim Harley**<sup>1</sup>

**Timothy P. Lillicrap**<sup>1</sup>

**David Silver**<sup>1</sup>

**Koray Kavukcuoglu**<sup>1</sup>

<sup>1</sup> Google DeepMind

<sup>2</sup> Montreal Institute for Learning Algorithms (MILA), University of Montreal

VMNIH@GOOGLE.COM

ADRIAP@GOOGLE.COM

MIRZAMOM@IRO.UMONTREAL.CA

GRAVESA@GOOGLE.COM

THARLEY@GOOGLE.COM

COUNTZERO@GOOGLE.COM

DAVIDSILVER@GOOGLE.COM

KORAYK@GOOGLE.COM

**A2C = Synchronous, Deterministic A3C**

## Other papers that help understand this work

---

### Policy Gradient Methods for Reinforcement Learning with Function Approximation

---

Fundamental math proofs

Richard S. Sutton, David McAllester, Satinder Singh, Yishay Mansour  
AT&T Labs – Research, 180 Park Avenue, Florham Park, NJ 07932

Parallel DQN

---

### Massively Parallel Methods for Deep Reinforcement Learning

---

Arun Nair, Praveen Srinivasan, Sam Blackwell, Cagdas Alcicek, Rory Fearon, Alessandro De Maria, Vedavyas Panneershelvam, Mustafa Suleyman, Charles Beattie, Stig Petersen, Shane Legg, Volodymyr Mnih, Koray Kavukcuoglu, David Silver

{ARUNSNIR, PRAV, BLACKWELLS, CAGDASALCICEK, RORYF, ADEMARIA, DARTHVEDA, MUSTAFASUL, CBEATTIE, SVP, LEGG, VMNIH, KORAYK, DAVIDSILVER @GOOGLE.COM }

Google DeepMind, London

---

### Dueling Network Architectures for Deep Reinforcement Learning

---

Ziyu Wang  
Tom Schaul  
Matteo Hessel  
Hado van Hasselt  
Marc Lanctot  
Nando de Freitas

Google DeepMind, London, UK

ZIYU@GOOGLE.COM  
SCHAUL@GOOGLE.COM  
MTTHSS@GOOGLE.COM  
HADO@GOOGLE.COM  
LANCTOT@GOOGLE.COM  
NANDODEFREITAS@GMAIL.COM

Advanced  
architecture

## What is A2C and A3C? Why asynchronous?

### Motivation

- To utilize more intuitions (tune *inner critics*, e.g. *self-driving cars: if the car turns left (policy), does it still in the correct route? is it closer to the destination? (critics, state value)* )
- Training data is too large to be trained efficiently (model is too large)
- Intensive uses of hardwares (GPUs)
- Make sure reducing correlations in training data by parallel actor-learners (replace experience replay buffer)

## Recall... Basics Value-based function approximation

Accumulative Return:

$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}$$

$$Q^{\pi}(s, a) = \mathbb{E}[R_t | s_t = s, a]$$

$$V^{\pi}(s) = \mathbb{E}[R_t | s_t = s]$$

One-step Q learning

$$L_i(\theta_i) = \mathbb{E} \left( r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right)^2$$

where  $s'$  is the state encountered after state  $s$ .

# Recall... Basics Policy-based RL

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

learning rate factor      reinforcement baseline

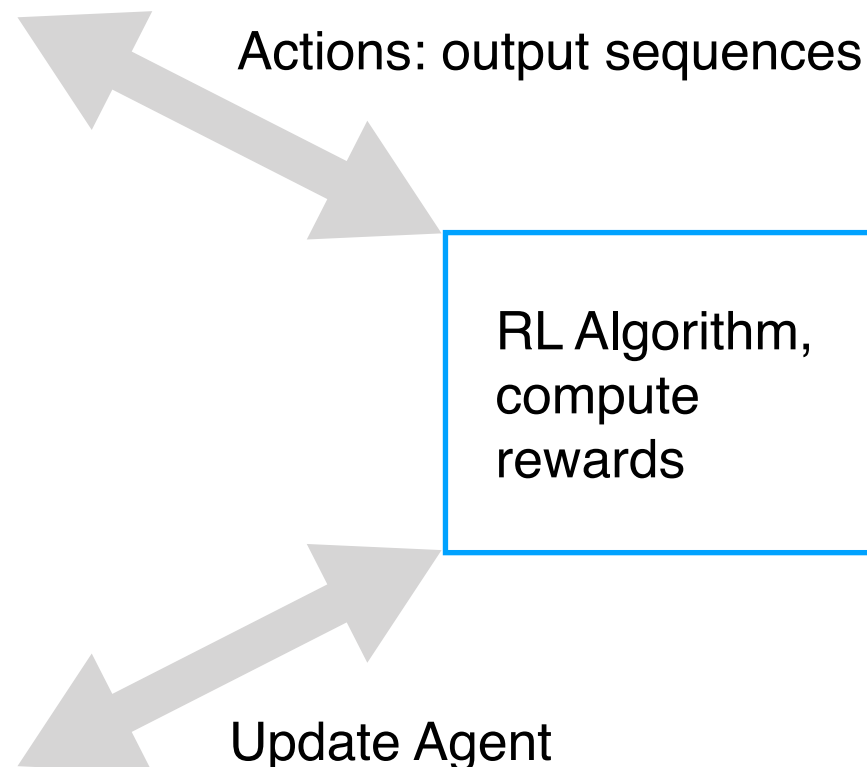
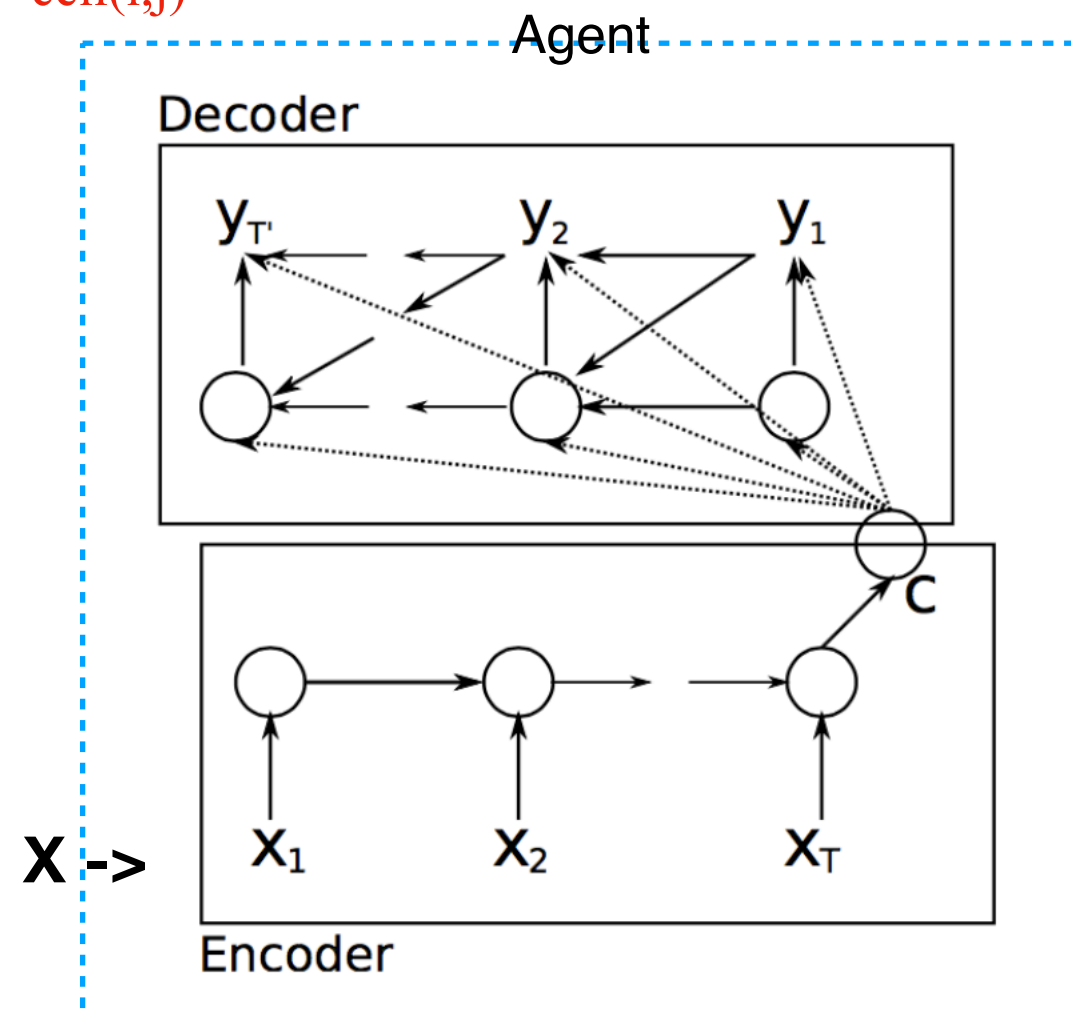
$$\Delta w_{ij} = \alpha_{ij} (r - b_{ij}) e_{ij},$$

changes of weight in each cell(i,j)

loss function

$$e_{ij} = \partial \ln g_i / \partial w_{ij}$$

characteristic eligibility of w(i,j)



Goal: learn a set of parameters that generate the best output (maximum rewards)

Could be plugged into any reward function, any networks



## Definitions for A3C (A2C)

(Baseline:)  $b_t(s_t) \approx V^\pi(s_t)$

*Asynchronous:* Parallel architectures (multi-thread execution) to improve training efficiency

*Advantage:*  $A(a_t, s_t) = Q(a_t, s_t) - V(s_t)$

$$A(s, a) = \underbrace{Q(s, a)}_{\substack{\text{estimation of} \\ \text{rewards}}} - \underbrace{V(s)}_{\substack{\text{estimation of} \\ \text{baseline}}}$$

q value for action a  
in state s
average  
value  
of that  
state

Change gradients update direction by  $A(s, a)$ :

- if  $A(s, a) > 0$ : extra rewards; update following the gradients direction
- if  $A(s, a) < 0$ : current selection is not wise; follow the opposite direction

*Actor-Critics:*

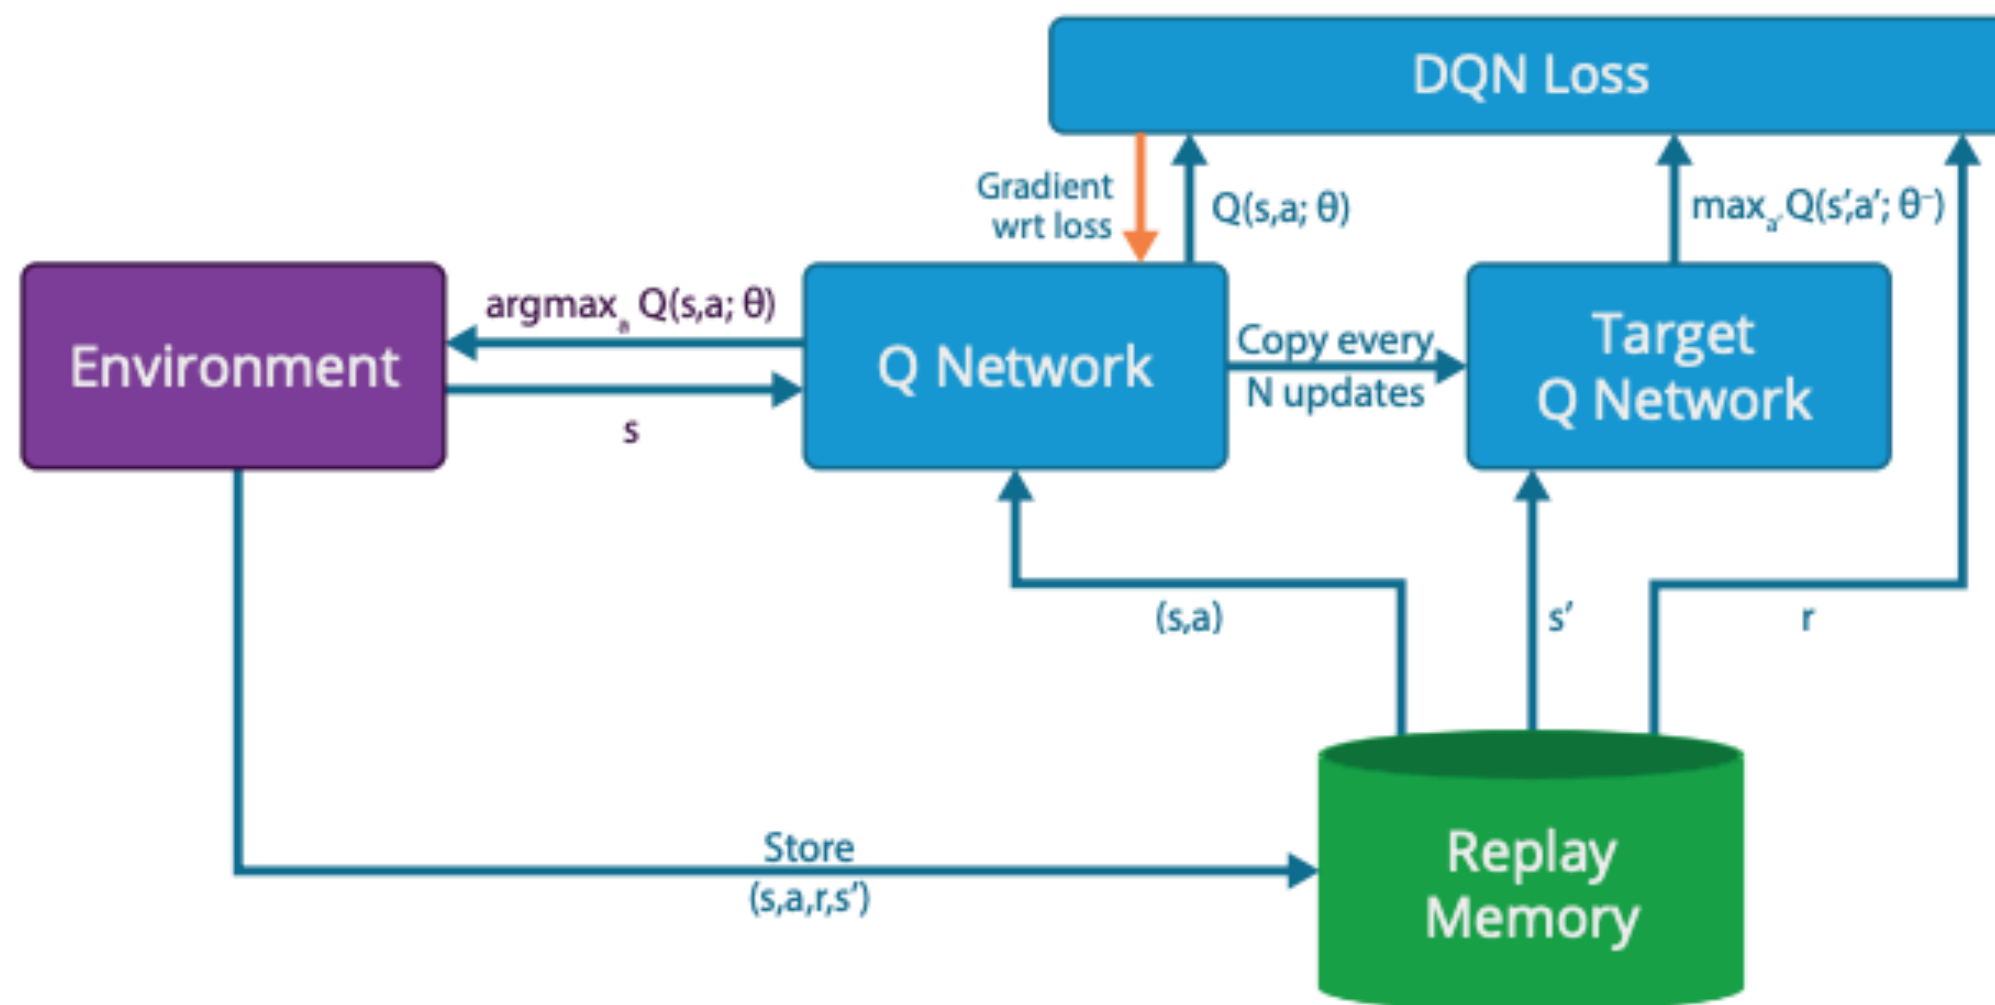
$\pi(s, a, \Theta)$ : Actor, policy function for telling agents how to act;

$Q(s, a, w)$ : Critics, value function for measuring how good the action is

## Definitions for A3C (A2C)

*Asynchronous*

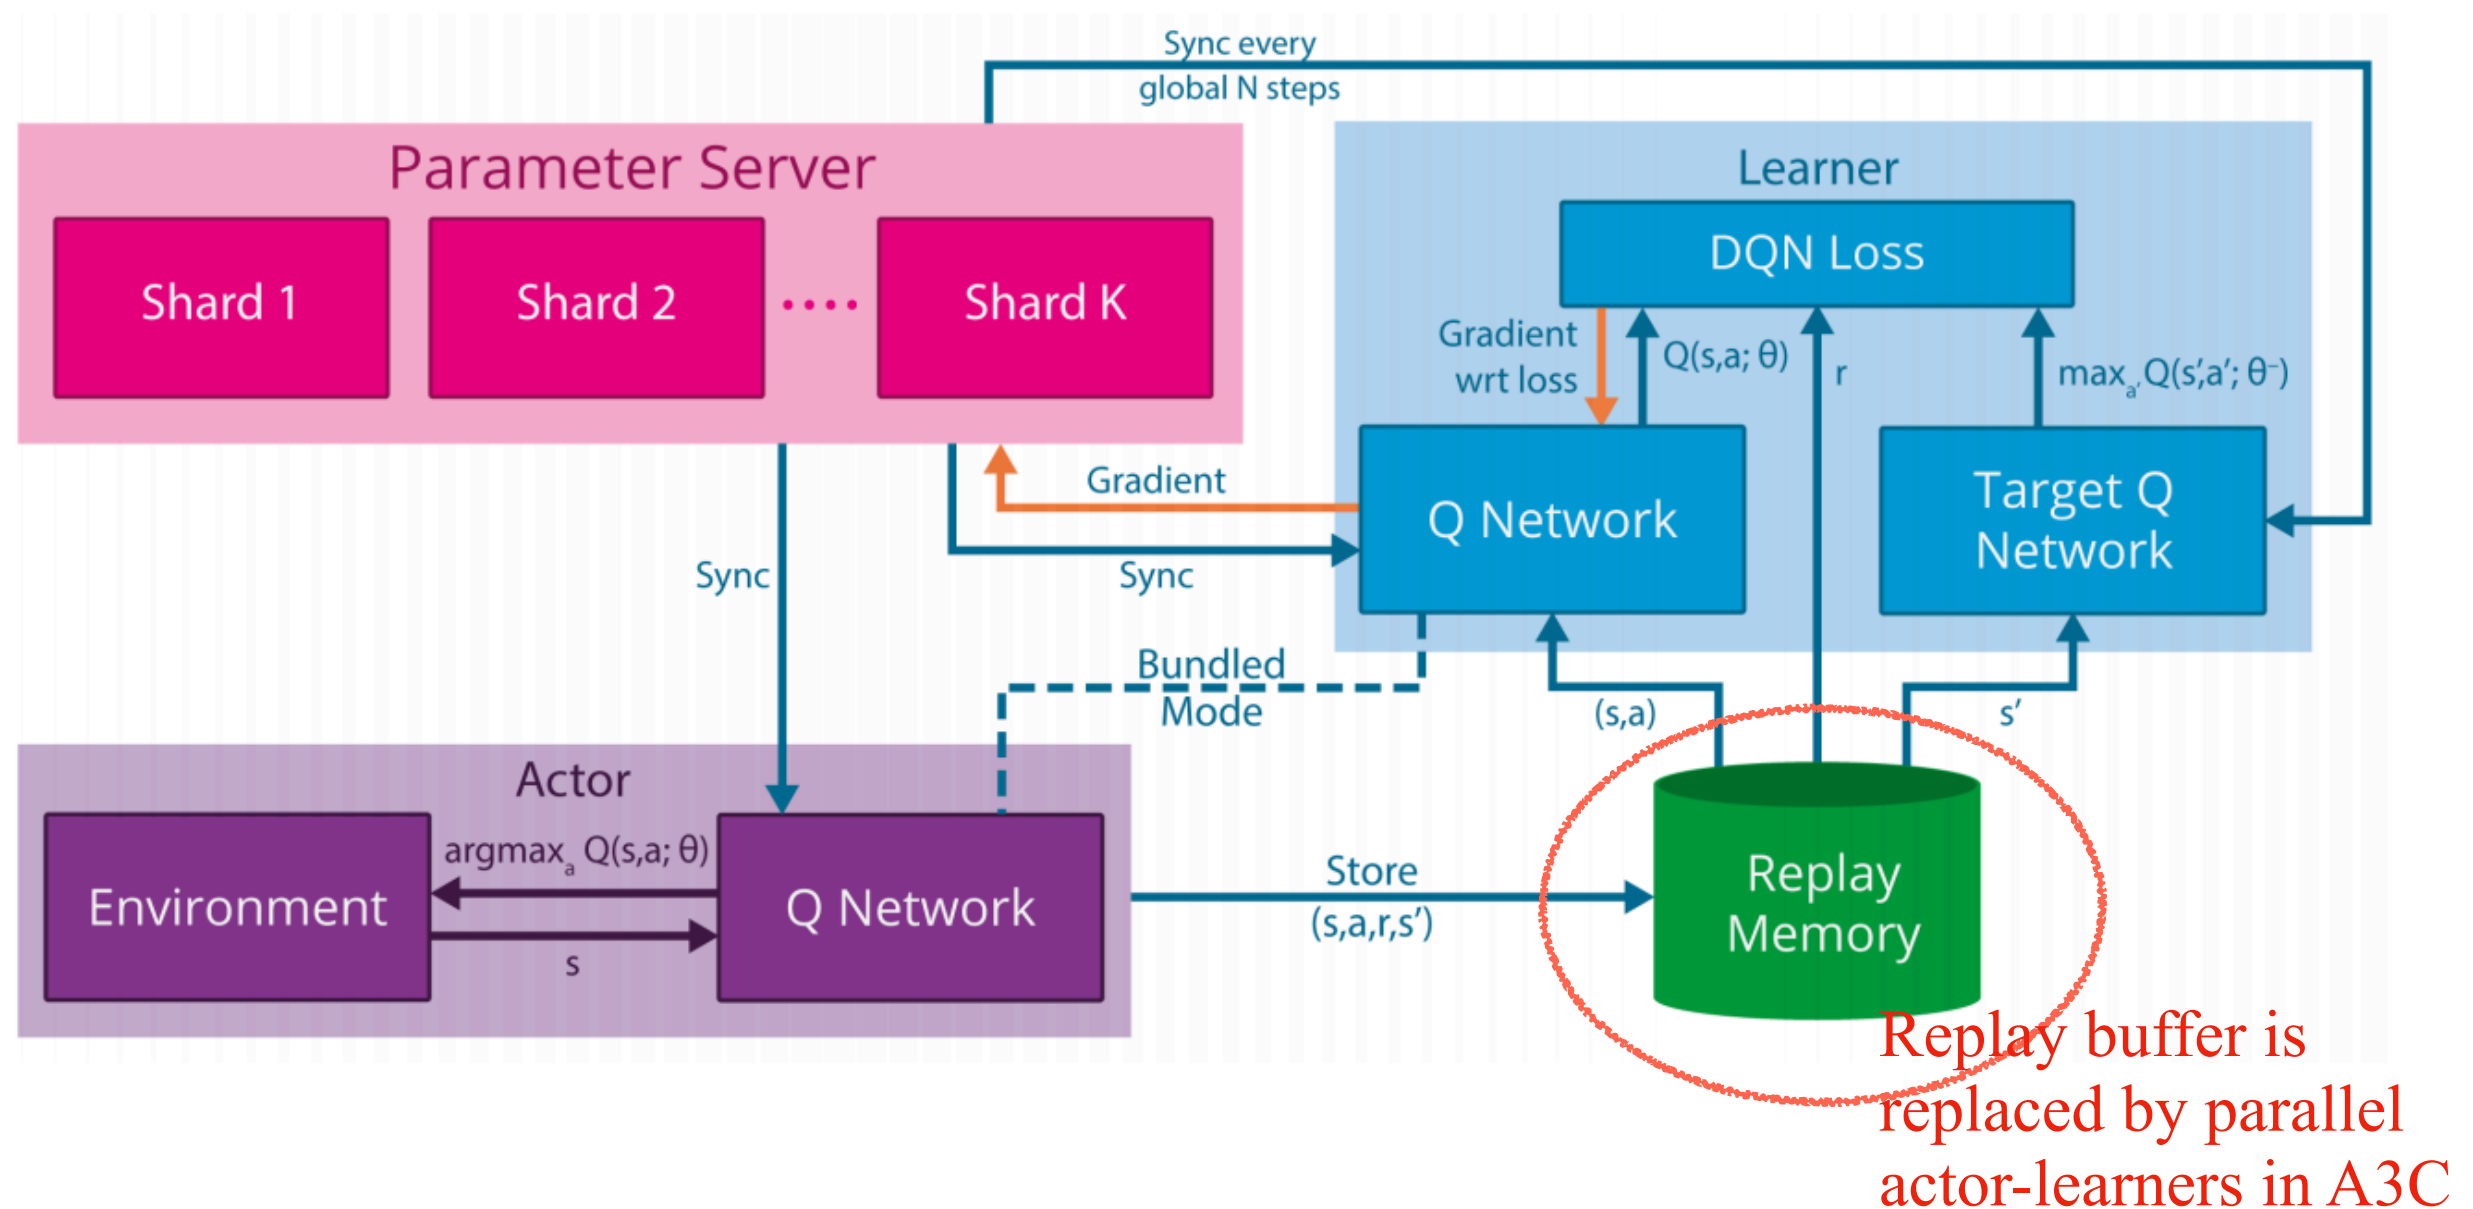
Take DQN as example, tradition DQN architecture:



## Definitions for A3C (A2C)

*Asynchronous*


Take DQN as example, **parallel** DQN architecture proposed in previous work:



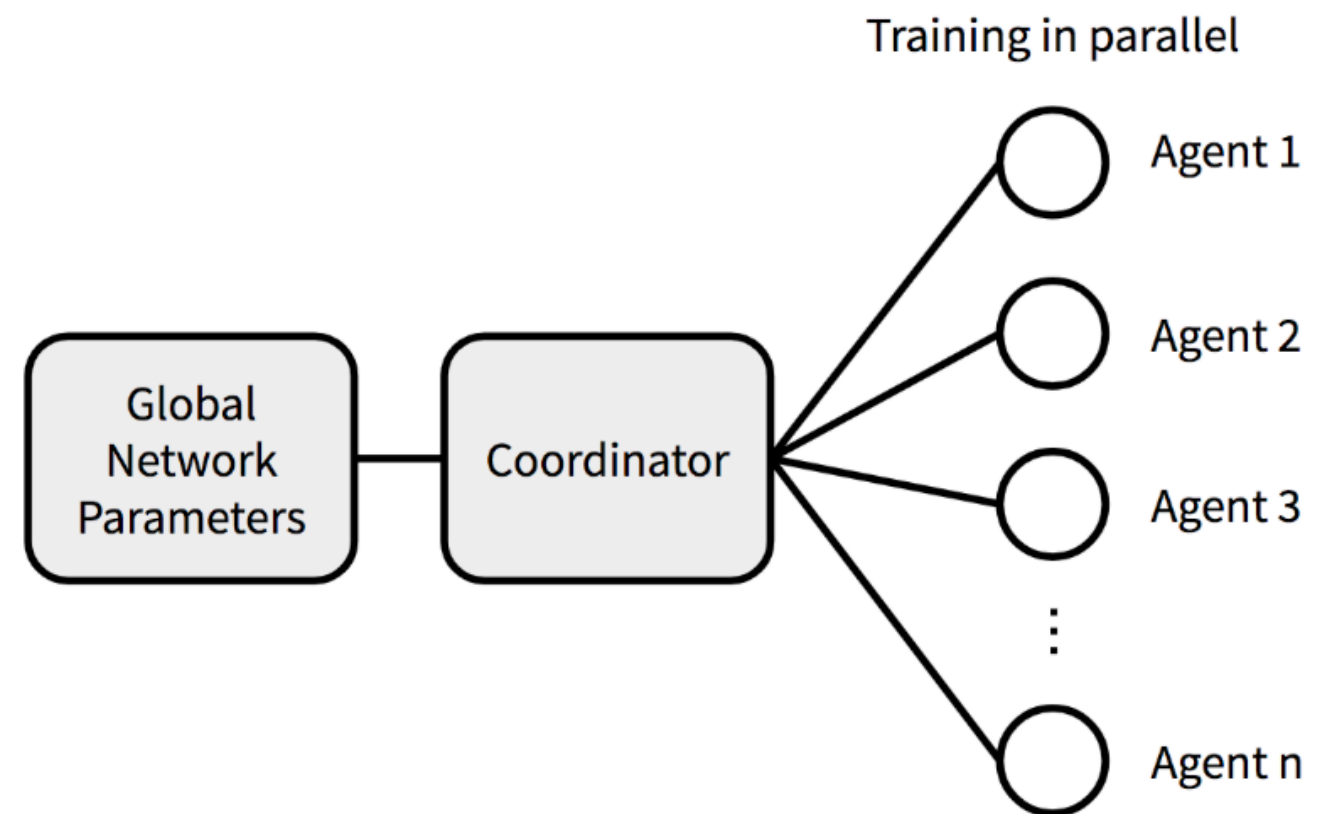
## Definitions for A3C (A2C)

*Asynchronous*

A3C VS A2C:



**A3C (Async)**



**A2C (Sync)**

Source: <https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html#a2c>

## Definitions for A3C (A2C)

*Asynchronous*

One-step Q-learning

---

**Algorithm 1** Asynchronous one-step Q-learning - pseudocode for each actor-learner thread.

---

```

// Assume global shared  $\theta$ ,  $\theta^-$ , and counter  $T = 0$ .
Initialize thread step counter  $t \leftarrow 0$ 
Initialize target network weights  $\theta^- \leftarrow \theta$ 
Initialize network gradients  $d\theta \leftarrow 0$ 
Get initial state  $s$ 
repeat
    Take action  $a$  with  $\epsilon$ -greedy policy based on  $Q(s, a; \theta)$ 
    Receive new state  $s'$  and reward  $r$ 
     $y = \begin{cases} r & \text{for terminal } s' \\ r + \gamma \max_{a'} Q(s', a'; \theta^-) & \text{for non-terminal } s' \end{cases}$ 
    Accumulate gradients wrt  $\theta$ :  $d\theta \leftarrow d\theta + \frac{\partial (y - Q(s, a; \theta))^2}{\partial \theta}$ 
     $s = s'$ 
     $T \leftarrow T + 1$  and  $t \leftarrow t + 1$ 
    if  $T \bmod I_{target} == 0$  then
        Update the target network  $\theta^- \leftarrow \theta$ 
    end if
    if  $t \bmod I_{AsyncUpdate} == 0$  or  $s$  is terminal then
        Perform asynchronous update of  $\theta$  using  $d\theta$ .
        Clear gradients  $d\theta \leftarrow 0$ .
    end if
until  $T > T_{max}$ 

```

---

## Definitions for A3C (A2C)

*Advantage:*  $\nabla_{\theta'} \log \pi(a_t|s_t; \theta') A(s_t, a_t; \theta, \theta_v)$

$A(s_t, a_t; \theta, \theta_v)$  is the estimate of  $\sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$

---

### Algorithm S3 Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

---

*// Assume global shared parameter vectors  $\theta$  and  $\theta_v$  and global shared counter  $T = 0$*

*// Assume thread-specific parameter vectors  $\theta'$  and  $\theta'_v$*

Initialize thread step counter  $t \leftarrow 1$

**repeat**

Reset gradients:  $d\theta \leftarrow 0$  and  $d\theta_v \leftarrow 0$ .

Synchronize thread-specific parameters  $\theta' = \theta$  and  $\theta'_v = \theta_v$

$t_{start} = t$

Get state  $s_t$

**repeat**

Perform  $a_t$  according to policy  $\pi(a_t|s_t; \theta')$

Receive reward  $r_t$  and new state  $s_{t+1}$

$t \leftarrow t + 1$

$T \leftarrow T + 1$

**until** terminal  $s_t$  **or**  $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

**for**  $i \in \{t-1, \dots, t_{start}\}$  **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

**end for**

Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .

**until**  $T > T_{max}$

---



# Definitions for A3C (A2C)

*Advantage:*  $\nabla_{\theta'} \log \pi(a_t|s_t; \theta') A(s_t, a_t; \theta, \theta_v)$

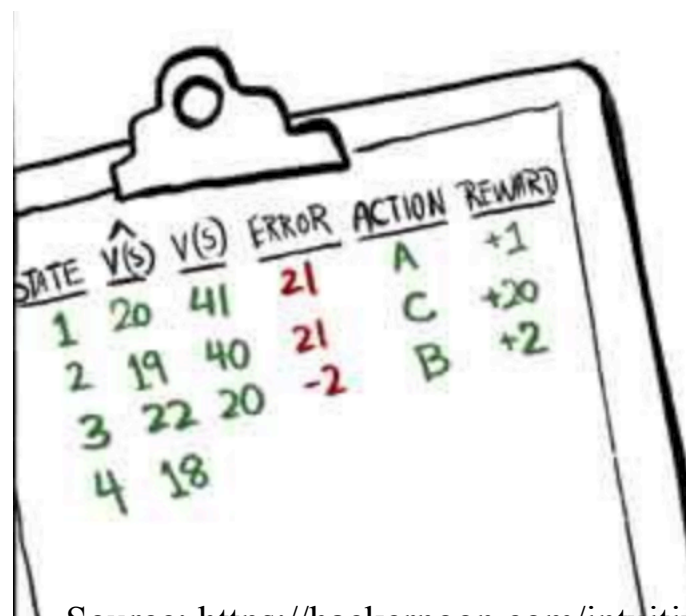
$A(s_t, a_t; \theta, \theta_v)$  is the estimate of  $\sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$   
 why not Q?

Recall:  $A(a_t, s_t) = Q(a_t, s_t) - V(s_t)$ ,  
 estimation of rewards      estimation of baseline

Of course you could compute Q (as in Dueling DQN), but ...

Recall:  $Q(s_t, a_t) = r_t + \sum_{k=0}^{\infty} \gamma^k \max Q_{t+k}(s_{t+k}, a_{t+k})$  Looking ahead of which action to take that leads to the best reward estimates (n-step Q)

Instead of looking ahead, A3C use its own learned critics (td-error) and prediction against the critic



STATE	$\hat{V}(s)$	$V(s)$	ERROR	ACTION	REWARD
1	20	41	21	A	+1
2	19	40	21	C	+20
3	22	20	-2	B	+2
4	18				

$$\sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$$

How many rewards actually are earned?

If continue, how many rewards one could earn? (Prediction)

Current estimate (looking back each row)

- How am I performing so far (MONTE CARLO)?
- How will I perform if I continue this way?
- How did I and will I perform compared to my expectation (estimation)?

## Definitions for A3C (A2C)

*Advantage:* TD - Error (Example of Richard Sutton driving home)

<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

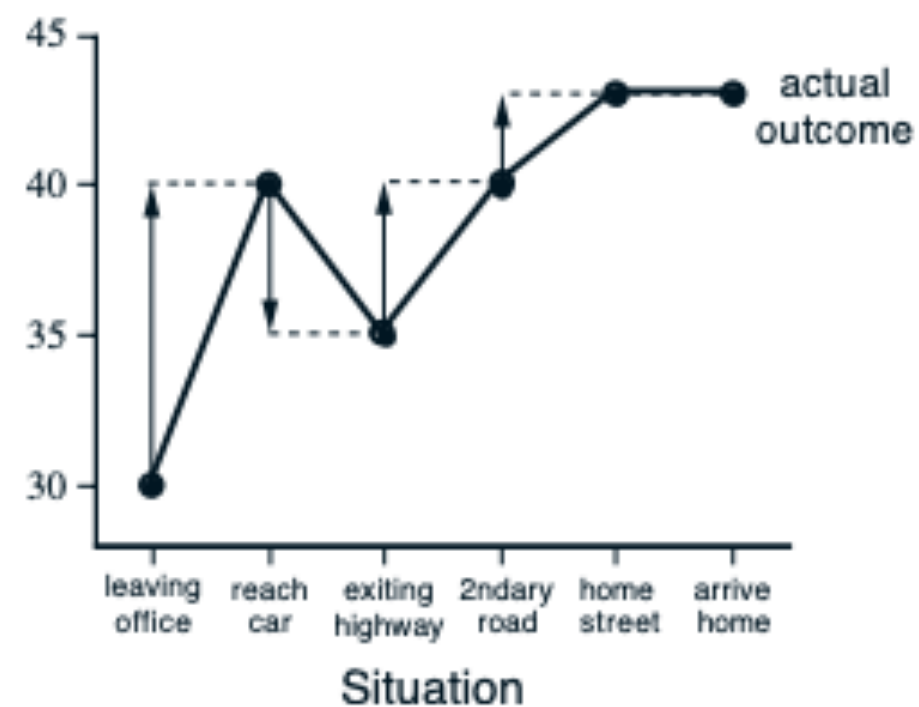



Figure 6.1: Changes recommended in the driving home example by Monte Carlo methods (left) and TD methods (right).



## Definitions for A3C (A2C)

*Actor-Critics:*

```
for  $i \in \{t-1, \dots, t_{start}\}$  do  
     $R \leftarrow r_i + \gamma R$   
    Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$   
    Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$   
end for
```

$$\nabla_{\theta'} \log \pi(a_t|s_t; \theta')(R_t - V(s_t; \theta_v)) + \beta \nabla_{\theta'} H(\pi(s_t; \theta'))$$

Entropy term that helps  
exploration

## Experiments and Results

Method	Training Time	Mean	Median
DQN	8 days on GPU	121.9%	47.5%
Gorila	4 days, 100 machines	215.2%	71.3%
D-DQN	8 days on GPU	332.9%	110.9%
Dueling D-DQN	8 days on GPU	343.8%	117.1%
Prioritized DQN	8 days on GPU	463.6%	127.6%
A3C, FF	1 day on CPU	344.1%	68.2%
A3C, FF	4 days on CPU	496.8%	116.6%
A3C, LSTM	4 days on CPU	623.0%	112.6%

*Table 1.* Mean and median human-normalized scores on 57 Atari games using the human starts evaluation metric. Supplementary Table [SS3](#) shows the raw scores for all games.

Asynchronous Methods for Deep Reinforcement Learning: TORCS

Sonic (not from this work)

## Take-away

- It doesn't mean experience replay is not useful even though A3C didn't use it; just more expensive
- A3C is able to handle complicated tasks, e.g. car racing, it combines both forward view (prediction) and backward view (eligible traceability)
- Actor-critic will be the mainstream of future reinforcement learning!

## References

- Mnih, Volodymyr, et al. "Asynchronous methods for deep reinforcement learning." *International conference on machine learning*. 2016.
- Sutton, Richard S., et al. "Policy gradient methods for reinforcement learning with function approximation." *Advances in neural information processing systems*. 2000.
- Nair, Arun, et al. "Massively parallel methods for deep reinforcement learning." *arXiv preprint arXiv:1507.04296* (2015). Make sure reducing correlations in training data by parallel actor-learners (replace experience replay buffer)
- Wang, Ziyu, et al. "Dueling network architectures for deep reinforcement learning." *arXiv preprint arXiv:1511.06581* (2015).
- <https://hackernoon.com/intuitive-rl-intro-to-advantage-actor-critic-a2c-4ff545978752>
- <https://medium.freecodecamp.org/an-intro-to-advantage-actor-critic-methods-lets-play-sonic-the-hedgehog-86d6240171d>
- Source: <https://cs.wmich.edu/~trenary/files/cs5300/RLBook/node66.html>
- Sutton, Richard S., and Andrew G. Barto. *Introduction to reinforcement learning*. Vol. 135. Cambridge: MIT press, 1998.