

# Reinforcement Learning For Sequence Learning in Neural Network

September 28, 2018

Yanjun Gao

PSU NLP Lab

- Problems
- RL Framework in Neural Network
- Sequence Learning with RL Framework

# Research Highlight - Reinforcement Learning in semantic representation

- Why using RL in semantic representation?
  - Most semantic representations are supervised learning. What if the model gets a sentence it has never seen before?
  - Sometimes the approaches are highly relied on the evaluation: how to address this with a good performance? Does one representation fit to all tasks?
  - RL is a hot topic; -> more attentions and publications!
- Why using RL in semantic representation is hard?
  - RL needs definitions of “agent, ” “states” , “actions” , “rewards” , “policy” ; it is a really hard process to formulate a semantic representation problem to RL problem; semantic representation is really different from dialogue modeling, auto-driving, robotics,etc.
  - Tradeoff between exploration and exploitation
  - Training data: where to get the training data is always a big quesiton;
  - Accuracy and performance: does RL always improve performance?
  - Other difficulties: optimizations, etc.

# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

learning rate factor      reinforcement baseline

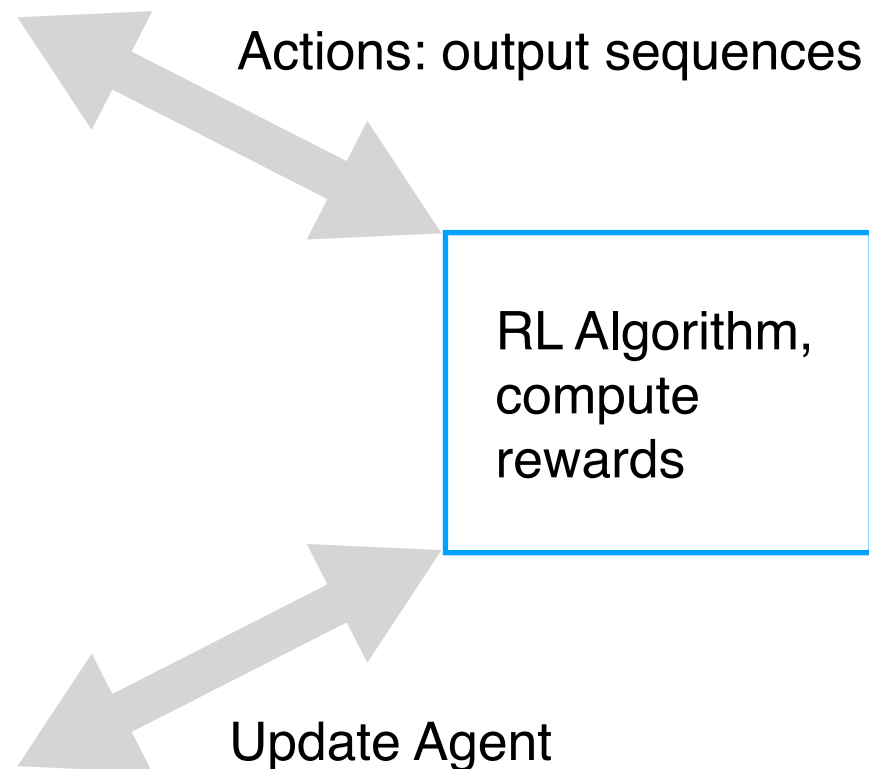
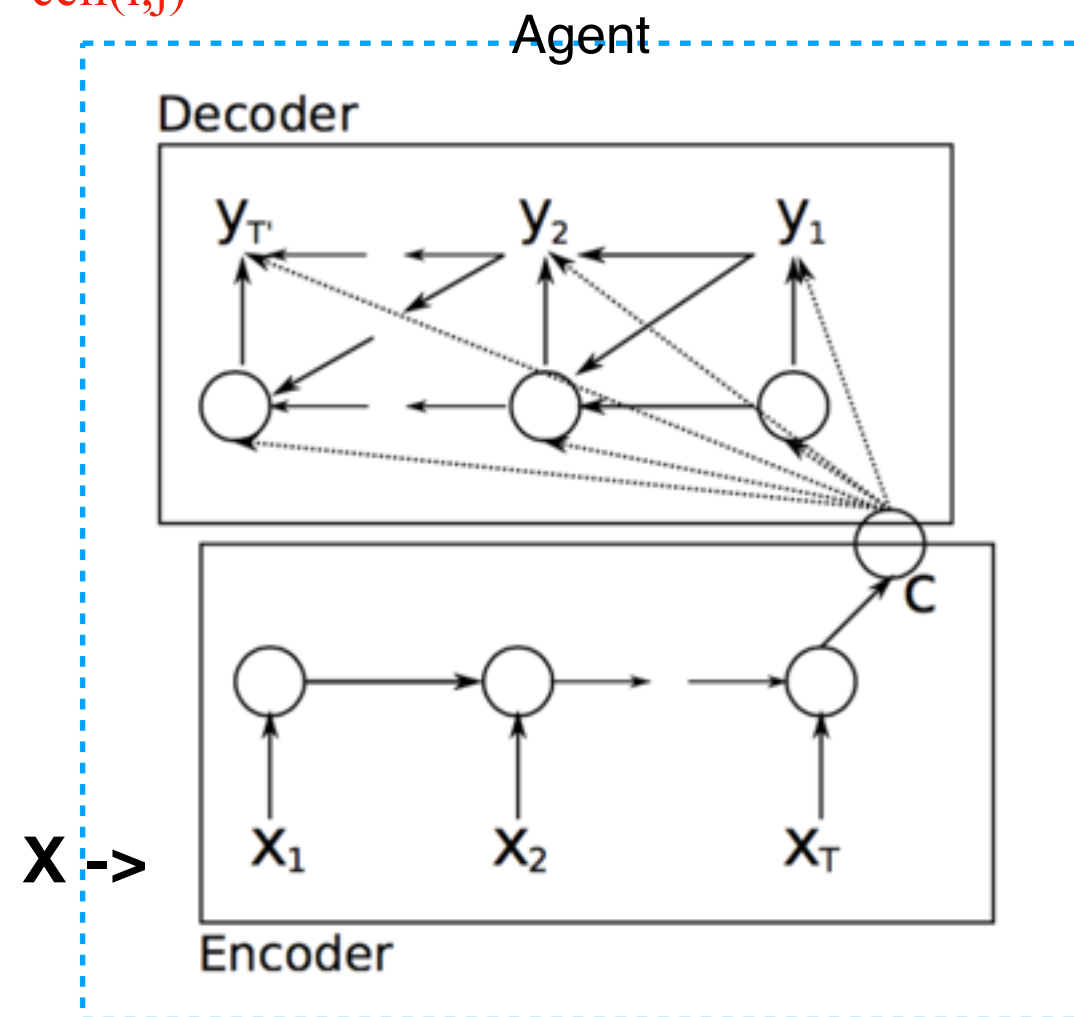
$$\Delta w_{ij} = \alpha_{ij} (r - b_{ij}) e_{ij},$$

changes of weight in each cell(i,j)

loss function

$$e_{ij} = \partial \ln g_i / \partial w_{ij}$$

characteristic eligibility of w(i,j)



Goal: learn a set of parameters that generate the best output (maximum rewards)

Could be plugged into any reward function, any networks

# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

$y_i$ : output of  $i$ th unit of the network;

$\mathbf{x}_i$ , pattern of input to that unit  $i$ .

The output  $y_i$  is drawn from the distribution depending on  $\mathbf{x}_i$  and weights  $\mathbf{w}_i$ .

$\mathbf{W}$ : weight matrix of all  $\mathbf{w}_i$ ;

$\mathbf{w}_i$ : collection of all parameters on  $i$ \_th unit.

$$\text{for each } i \text{ let } g_i(\xi, \mathbf{w}^i, \mathbf{x}^i) = \Pr\{y_i = \xi | \mathbf{w}^i, \mathbf{x}^i\}.$$

# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

Then we define a *stochastic semilinear unit* to be one whose output  $y_i$  is drawn from some given probability distribution whose mass function has single parameter  $p_i$ :

$$p_i = f_i(s_i), \tag{1}$$

$f_i$ : differentiable squashing function,  
e.g. logistic function

$$s_i = \mathbf{w}^{iT} \mathbf{x}^i = \sum_j w_{ij} x_j, \tag{2}$$

# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

When  $y_i$  follows Bernoulli distribution, we define Bernoulli semilinear unit with parameter  $p_i$ , where all possible output values are 0 and 1. Thus:

$$g_i(\xi, \mathbf{w}^i, \mathbf{x}^i) = \begin{cases} 1 - p_i & \text{if } \xi = 0 \\ p_i & \text{if } \xi = 1, \end{cases}$$

Alternatively,

$$y_i = \begin{cases} 1 & \text{if } \sum_j w_{ij} x_j + \eta > 0 \\ 0 & \text{otherwise,} \end{cases}$$

# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

Given a parameter to be adapted  $p_i = \Pr(y_i=1)$ , follows Bernouli distribution

$$g_i(y_i, p_i) = \begin{cases} 1 - p_i & \text{if } y_i = 0 \\ p_i & \text{if } y_i = 1, \end{cases} \quad (4)$$

Define characteristic eligibility for parameter  $p_i$ :

$$\frac{\partial \ln g_i}{\partial p_i}(y_i, p_i) = \begin{cases} -\frac{1}{1-p_i} & \text{if } y_i = 0 \\ \frac{1}{p_i} & \text{if } y_i = 1 \end{cases} = \frac{y_i - p_i}{p_i(1 - p_i)}, \quad (5)$$

Define characteristic eligibility for weight  $w_{ij}$ :

$$\frac{\partial \ln g_i}{\partial w_{ij}}(y_i, \mathbf{w}^i, \mathbf{x}^i) = \frac{y_i - p_i}{p_i(1 - p_i)} f_i'(s_i) x_j, \quad (6)$$



# A RL framework for parametric neural network (Williams, 1992)

Williams, Ronald J. "Simple statistical gradient-following algorithms for connectionist reinforcement learning." *Machine learning* 8.3-4 (1992): 229-256.

Set  $f_i$  as logistic function,  $b_{ij} = 0$ ,

$$\frac{\partial \ln g_i}{\partial w_{ij}}(y_i, \mathbf{w}^i, \mathbf{x}^i) = (y_i - p_i)x_j. \quad (7)$$

Through some substitutions, we have the form:

$$\Delta w_{ij} = \alpha r (y_i - p_i)x_j,$$

Adopt “reinforcement comparison” in (Sutton 1984),  $\hat{r}$  is an adaptive estimate of upcoming reinforcement based on past experience:

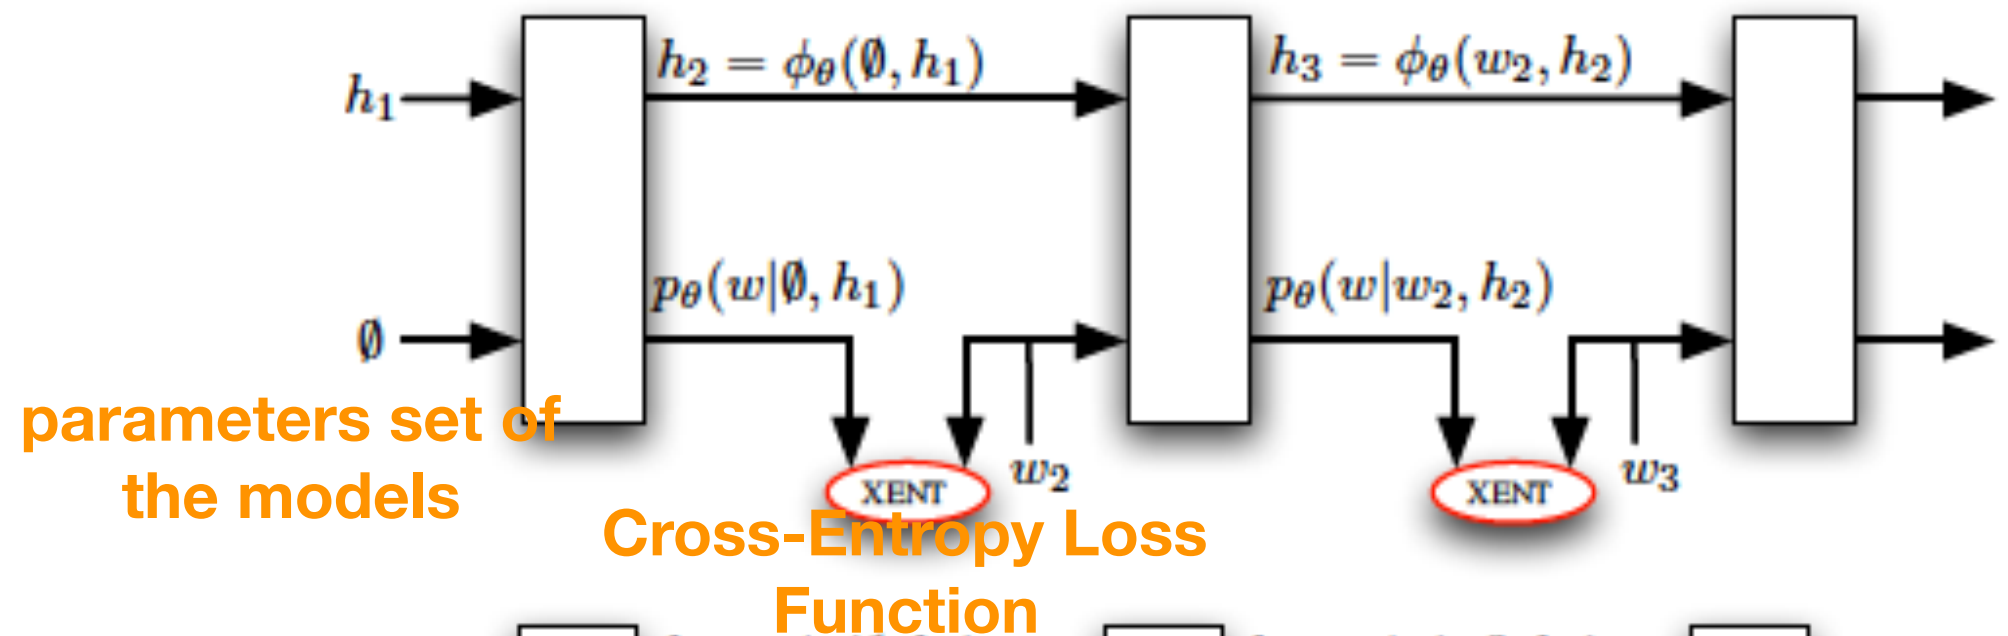
$$\Delta w_{ij} = \alpha (r - \bar{r})(y_i - p_i)x_j, \quad (9)$$

# Sequence Level Training with Recurrent Neural Networks


Ranzato, Marc'Aurelio, et al. "Sequence level training with recurrent neural networks." *arXiv preprint arXiv:1511.06732* (2015).

## Typical Elman RNN without RL

Training:



Testing (model prediction):



$$\mathbf{h}_{t+1} = \sigma(M_i \mathbf{1}(w_t) + M_h \mathbf{h}_t + M_c \mathbf{c}_t), \quad (3)$$

$$\mathbf{o}_{t+1} = M_o \mathbf{h}_{t+1}, \quad (4)$$

$$w_{t+1} \sim \text{softmax}(\mathbf{o}_{t+1}), \quad (5)$$

# Sequence Level Training with Recurrent Neural Networks

Ranzato, Marc'Aurelio, et al. "Sequence level training with recurrent neural networks." *arXiv preprint arXiv:1511.06732* (2015).

## Mixed Incremental Cross-Entropy Reinforce (MIXER)

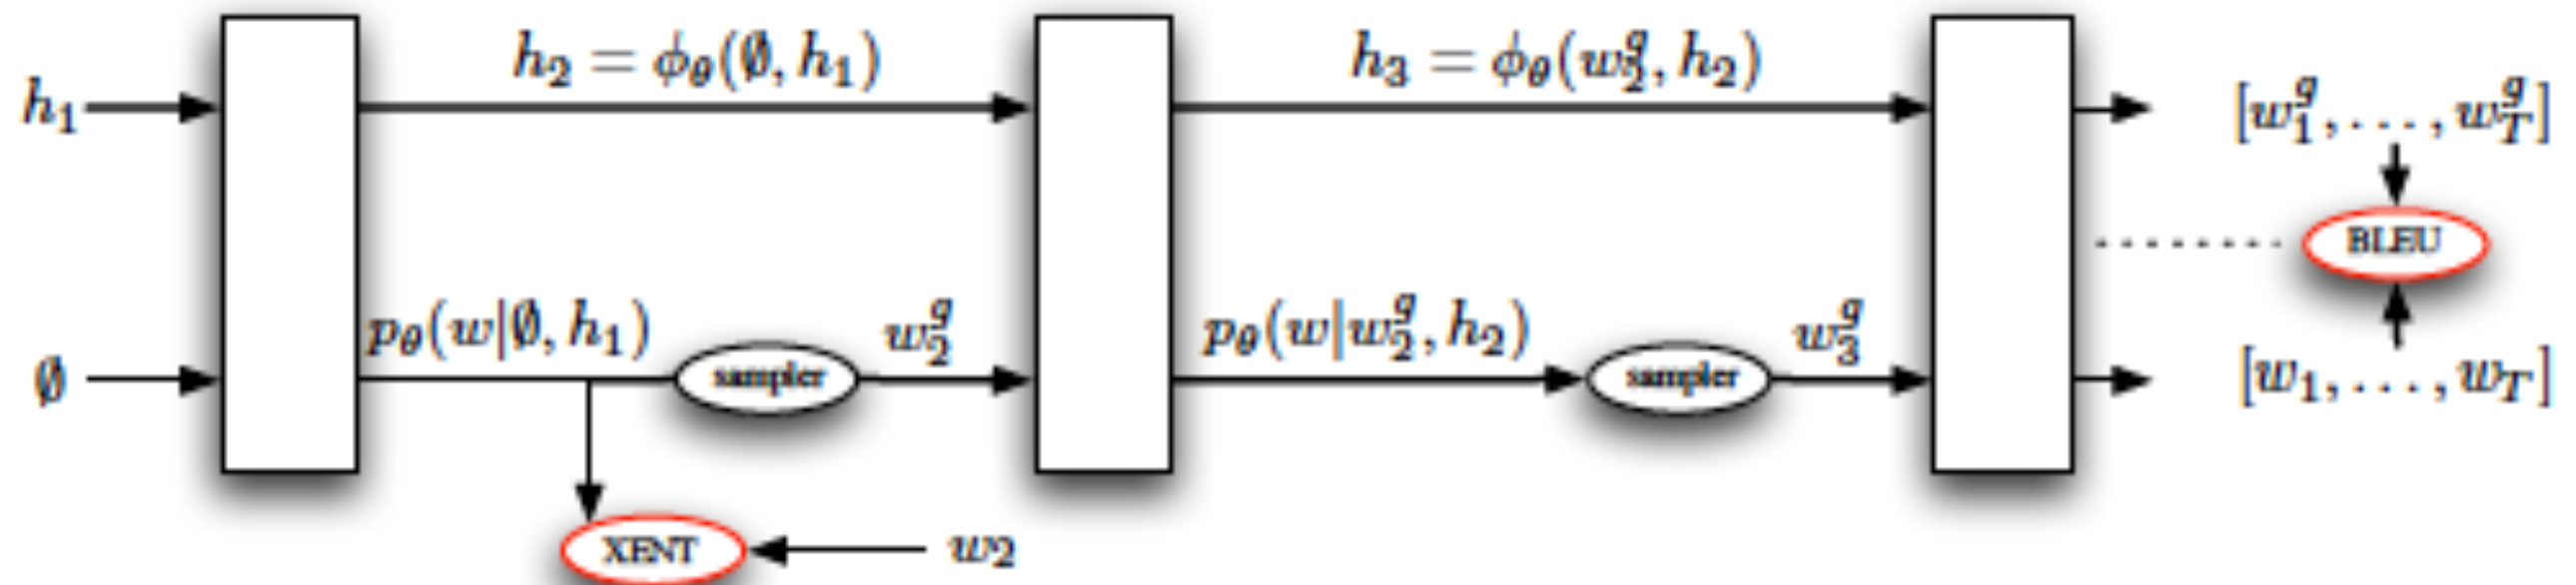


Figure 4: Illustration of MIXER. In the first  $s$  unrolling steps (here  $s = 1$ ), the network resembles a standard RNN trained by XENT. In the remaining steps, the input to each module is a sample from the distribution over words produced at the previous time step. Once the end of sentence is reached (or the maximum sequence length), a reward is computed, e.g., BLEU. REINFORCE is then used to back-propagate the gradients through the sequence of samplers. We employ an annealing schedule on  $s$ , starting with  $s$  equal to the maximum sequence length  $T$  and finishing with  $s = 1$ .

# Sequence Level Training with Recurrent Neural Networks

Ranzato, Marc'Aurelio, et al. "Sequence level training with recurrent neural networks." *arXiv preprint arXiv:1511.06732* (2015).

## Mixed Incremental Cross-Entropy Reinforce (MIXER)

$$L_{\theta} = - \sum_{w_1^g, \dots, w_T^g} p_{\theta}(w_1^g, \dots, w_T^g) r(w_1^g, \dots, w_T^g) = -\mathbb{E}_{[w_1^g, \dots, w_T^g] \sim p_{\theta}} r(w_1^g, \dots, w_T^g), \quad (9)$$

$$\frac{\partial L_{\theta}}{\partial \theta} = \sum_t \frac{\partial L_{\theta}}{\partial \mathbf{o}_t} \frac{\partial \mathbf{o}_t}{\partial \theta} \quad (10)$$

$$\frac{\partial L_{\theta}}{\partial \mathbf{o}_t} = (r(w_1^g, \dots, w_T^g) - \bar{r}_{t+1}) (p_{\theta}(w_{t+1} | w_t^g, \mathbf{h}_{t+1}, \mathbf{c}_t) - \mathbf{1}(w_{t+1}^g)), \quad (11)$$

# Sequence Level Training with Recurrent Neural Networks

Ranzato, Marc'Aurelio, et al. "Sequence level training with recurrent neural networks." *arXiv preprint arXiv:1511.06732* (2015).

**Data:** a set of sequences with their corresponding context.

**Result:** RNN optimized for generation.

Initialize RNN at random and set  $N^{\text{XENT}}$ ,  $N^{\text{XE+R}}$  and  $\Delta$ ;

for  $s = T, 1, -\Delta$  do

    if  $s == T$  then

first (T-delta) steps, train XENT

        train RNN for  $N^{\text{XENT}}$  epochs using XENT only;

    else

        train RNN for  $N^{\text{XE+R}}$  epochs. Use XENT loss in the first  $s$  steps, and REINFORCE (sampling from the model) in the remaining  $T - s$  steps;

Rest of (T-s) steps, anneal output to a stable sequence

    end

end

Algorithm 1: MIXER pseudo-code.



# Sequence Level Training with Recurrent Neural Networks

Ranzato, Marc'Aurelio, et al. "Sequence level training with recurrent neural networks." *arXiv preprint arXiv:1511.06732* (2015).

<i>TASK</i>	<b>XENT</b>	<b>DAD</b>	<b>E2E</b>	<b>MIXER</b>
<i>summarization</i>	13.01	12.18	12.78	<b>16.22</b>
<i>translation</i>	17.74	20.12	17.77	<b>20.73</b>
<i>image captioning</i>	27.8	28.16	26.42	<b>29.16</b>




Figure 5: Left: BLEU-4 (translation and image captioning) and ROUGE-2 (summarization) scores using greedy generation. Right: Relative gains produced by DAD, E2E and MIXER on the three tasks. The relative gain is computed as the ratio between the score of a model over the score of the reference XENT model on the same task. The horizontal line indicates the performance of XENT.