

# Reinforcement Learning

By

Saptarashmi Bandyopadhyay

First Year PhD Student

Department of Computer Science and  
Engineering

Pennsylvania State University, University Park

# Outline

- Reinforcement Learning Book
  - Introduction (Chapter 1)
  - Multi-armed bandit problem (Chapter 2)
  - Finite Markov decision Process (Chapter 3)
- A Course in Machine Learning
  - Imitation Learning (Chapter 18)
- Journal paper on Human-level control through deep reinforcement learning

# Introduction

- Re-inforcement Learning- Goal directed learning from interaction with the environment
- Learning how to map situations to actions to maximize a numerical reward signal
- Trial-and-error search to find which actions yield the most reward
- Delayed reward
  - Short-term reward
  - Long-term reward

# Learning agents

- Able to sense state of the environment
- Able to take actions based on affecting the state
- Goal or goals related to the state of the environment
- Complete interactive goal-seeking agent
- Can be a component of a bigger behaving system

# Difference with Supervised Learning

- Supervised learning: learning from a trained set of labeled known samples
- System to generalize for situations outside training dataset
- But in interactive problems, impractical to get desired behavior that are both correct and reflects all situations of the agent's action space
- Learning from experience in reinforcement learning

# Difference with Unsupervised Learning

- Unsupervised learning: learning of structures hidden in unlabeled data
- Reinforcement learning(RL) maximizes the goal reward instead of trying to find hidden structures
- Uncovering the structure useful but does not address the core objective of RL

# Trade-off between exploration vs. exploitation

- Exploitation of experience in order to obtain reward
- Exploration to select better actions
- Neither exploration nor exploitation can be done solely without task failure

# Illustration in chess

- Move of an expert chess player
- Choice by planning possible actions and counter-responses in the long term
- immediate, intuitive assessment of whether particular positions and moves are more favorable



# Elements of Reinforcement Learning (RL)

- Policy defines the behavior of learning agent
- Reward signal defines the goal of RL problem
- Value function highlighting long-term desirability of the state space
- Model defining interaction with the environment


# Limitations of Reinforcement Learning

- Highly dependent on the idea of state as input to the policy, value function and input and output of the model
  - Current focus on decision-making
- Problems can be unstructured around estimation of value functions like evolutionary methods e.g. genetic algorithms
  - Multiple static policies interacting over a long time with separate instances of the environment

# Illustration in tic-tac-toe game

- Mini-max algorithm assumes the playing process of the opponent
- Dynamic programming can compute optimal solution but requires complete specification of the opponent
- Information often unavailable apriori but can be estimated from experience
- $V(S^t) \leftarrow V(S^t) + \alpha * [V(S^{t+1}) - V(S^t)]$ 
  - $S^t$ : State before greedy move
  - $S^{t+1}$ : State after greedy move

# State space representation of tic-tac-toe game



# Multi-armed bandits problem

- Non-associative, evaluative feedback problem
- Choice from  $k$  different options
- After each choice to maximize numerical reward from a stationary probability distribution based on selected action
- $q_*(a) = E[R_t \mid A_t = a]$
- $q_*(a)$  : value of an arbitrary action  $a$
- $A_t$  is action at time  $t$  and  $R_t$  is corresponding reward

# Action value methods

- Methods to estimate value of actions for action-selection decisions
- Averaging the rewards
- Greedy action selection method
  - $A_t = \operatorname{argmax} Q_t(a)$
  - Where  $Q_t(a)$  is estimated action at time  $t$

# 10-armed testbed

- Set of 2000 randomly generated k-armed bandit problems with  $k=10$
- Comparison to find how effective greedy and  $\epsilon$ -greedy selection methods are
- $\epsilon$ -greedy selection
  - greedy behavior mostly
  - but sometimes with a small probability  $\epsilon$  random selection from among all actions with equal probability, independent of the action-value estimates

# Incremental Implementation

- Computationally efficient method for average in estimating action
- $Q_{n+1} = Q_n + (1/n) * [R_n - Q_n]$ 
  - Given  $Q_n$  and the  $n$ th reward  $R_n$
- General formula is
- New estimate  $\leftarrow$  old estimate + step-size\*(target - old estimate)



# Non-stationary problem

- In stationary bandit problems, reward probabilities do not change over time
- In non-stationary problems, more weightage to recent rewards than past rewards
- Exponential recency weighted average

# Optimistic initial value

- Previous methods biased by initial action value estimate  $Q_1(a)$
- Sample-average methods, bias disappears on selection of all subjects at least once
- Bias permanent for constant  $\alpha$  methods
- Exploration by action value methods to find optimistic initial values.

# Upper-confidence-bound action selection

- $A_t = \operatorname{argmax}[Q_t(a) + c * \sqrt{\ln(t)/N_t(a)}]$ 
  - $N_t(a)$  it is the number of times,  $a$  has been selected prior to time  $t$
- Square root term denotes the uncertainty or variance in the estimate of  $a$ 's value
- Maximum argument is an upper bound on the possible true value of action  $a$ .

# Gradient bandit problem

- Relative preference of one action over another
- Determined by a soft-max distribution

# Contextual bandit

- Associative search
- Several k-armed bandit problems
- At each step, one of the problems chosen at random
- Policy associating each task with the best action
- Actions affecting the next situation and reward giving the full reinforcement learning problem

# Finite Markov Decision Process

- Formalizes the problem of reinforcement learning
- Formalizes sequential decision making
- Actions influence immediate rewards and future rewards

# Agent-environment interface

- Agent: Learner and decision maker
- Environment: the thing agents interact with
- The probabilities completely characterize the environment's dynamics
- $p(s', r | s, a) = \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$ 
  - Dynamics function of MDP
- Markov property: state has to include all aspects of the past agent-environment interaction

# Goals and rewards

- Reward hypothesis
- Goals: maximization of the expected value of the cumulative sum of an obtained scalar signal called reward



# Returns and episodes

- Return is some specific function on the reward sequence
- Objective to maximize return
- Episodes: when agent-environment interaction naturally breaks into subsequences
- All episodes can end in the same terminal state with different rewards for different results
- Episodic task: Tasks with such episodes

# Unified notation for episodic and continuing tasks

- Episodic tasks need extra notation
- Rather than one long sequence of time steps, series of episodes with finite sequence of time steps

# Policies and value functions

- Policy: mapping of states to probabilities of selecting each possible action
- Value function of a state  $s$  and policy  $\pi$ : expected return when starting in  $s$  and following  $\pi$  thereafter.
- State value function for policy  $\pi$
- Action value function for policy  $\pi$

# Optimal policies and optimal value functions

- Optimal policy: Policy better than or equal to

all other policies

- Optimal state-value function:

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- Optimal action-value function:

$$q_*(s,a) = \max_{\pi} q_{\pi}(s,a)$$

# Optimality and approximation

- Extremely high computation cost
- Constraint of available memory
- Approximation e.g. tesauro's backgammon game
- online nature of reinforcement learning
- Approximates optimal policies
  - To learn to make good decisions for frequent states
  - Less effort for infrequent states

# References

Reinforcement learning book by Richard S. Sutton and Andrew G. Barto, Second Edition, 2018

# Imitation learning

- Learning by demonstration
- Access to an expert
- Goal: to learn function  $f$  that maps data from  
environmental interaction to actions
- Close relationship with reinforcement learning
- Function  $f$  is called a policy

# Certain nuances

- Conversion of expert data to multi-class classification problems
- Goodness of the expert
- Goodness of the multi-class classification algorithm
- Upper bound of loss suffered by supervised imitation learning algorithm is a function of
  - Expert quality
  - Error rate of learned classifier



# Failure Analysis

- Challenge is recovery from failure
- Compounding error
- Expert can easily get zero loss
- Upper and lower bounds of loss pretty far apart

# Dataset aggregation

- Ideally generalistic imitation learning
- Practically impossible to get expert data from all possible combinations in the world
- Train model based on configurations of the expert ?!
- So, dagger (dataset aggregation algorithm)
- Using expert to generate dataset for recovery
- Direct contact with experts
- But less compounding errors

# Expensive algorithms as experts

- Game playing : computationally expensive semi-optimal behavior computed with brute force search during simulation
  - rather to use the search during training to learn a fast policy mimicking the search
  - then tested
- discrete optimizers: computing shortest paths offline as training data and testing it

# Structured prediction via Imitation Learning

- Structured prediction can be solved via imitation learning using sequence labeling
- e.g. 'Monsters eat tasty bunnies'  
noun      verb adjective noun
- Policy to learn noun then verb then adjective then noun

# Reference

A Course in Machine Learning by Haul

Daumé III

<http://ciml.info/>

# Human-level control through deep reinforcement learning

- 19 authors from Google DeepMind, London
- 3 equal contribution authors Volodymyr Mnih, Koray Kavukcuoglu, David Silver with only one first author
- 3735 citations

# Background

- RL theory gives a normative account
- Psychological and neuro-scientific perspectives on animal behavior
- Efficient representation of environment from high dimensional sensory inputs
- General past experience with new situations

# Problems with reinforcement learning agents

It is restricted to domains with

- Hand-crafted useful features
- Full observed, low dimensional state spaces



# Work done

- Deep Q network, a novel artificial agent
- Learns successful policies from high dimensional sensory inputs using end to end reinforcement learning
- Testing on Atari 2600 games
- Pixels and game scores as inputs
- Result
  - Better performance than all previous algorithms
  - Comparable level with a professional human game tester over a set of 49 games

# Atari 2600

- Home video game console developed in 1977
- Use of general purpose CPUs in game console hardware
- Game code distributed through cartridges
- Arcade games like PACMAN ported to console
- Simple hardware limiting the game complexity
- Obtainable short-term progress in learning, modeling, and planning

# Architecture used

- Deep convolutional neural network
- Exploits local spatial correlation present in images
- Robust to transformations such as changes of viewpoints or scales
- Approximates optimal action value function

# Optimal action-value function

$$Q^*(s,a) = \max_{\pi} \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi],$$

maximum sum of rewards  $r_t$  discounted by  $\gamma$  at each time-step  $t$ , obtained by a behaviour policy  $\pi = P(a|s)$ , after an observation ( $s$ ) and executing an action ( $a$ )

# Reasons for instability of reinforcement learning

- Unstable or diverges when non linear function approximator represents the action value function
- Reasons:-
  - Correlations in the sequence of observations
  - Small updates to  $Q$  may change the policy and the data distribution
  - Correlations between the action values and target values

# Solution: Novel variant of Q Learning

- Experience replay randomizes over the data
  - Reducing correlation in observation sequence
  - Smoothing changes in data distribution
- Iterative update of action values towards target values that are periodically updated, reducing correlation with the targets

# Training algorithm

- Parameterizing an objective value function using deep neural network
- Action replay done by storing agent's experiences  $e_t = (s_t, a_t, r_t, s_{t+1})$  at each time step in a dataset

# Loss function

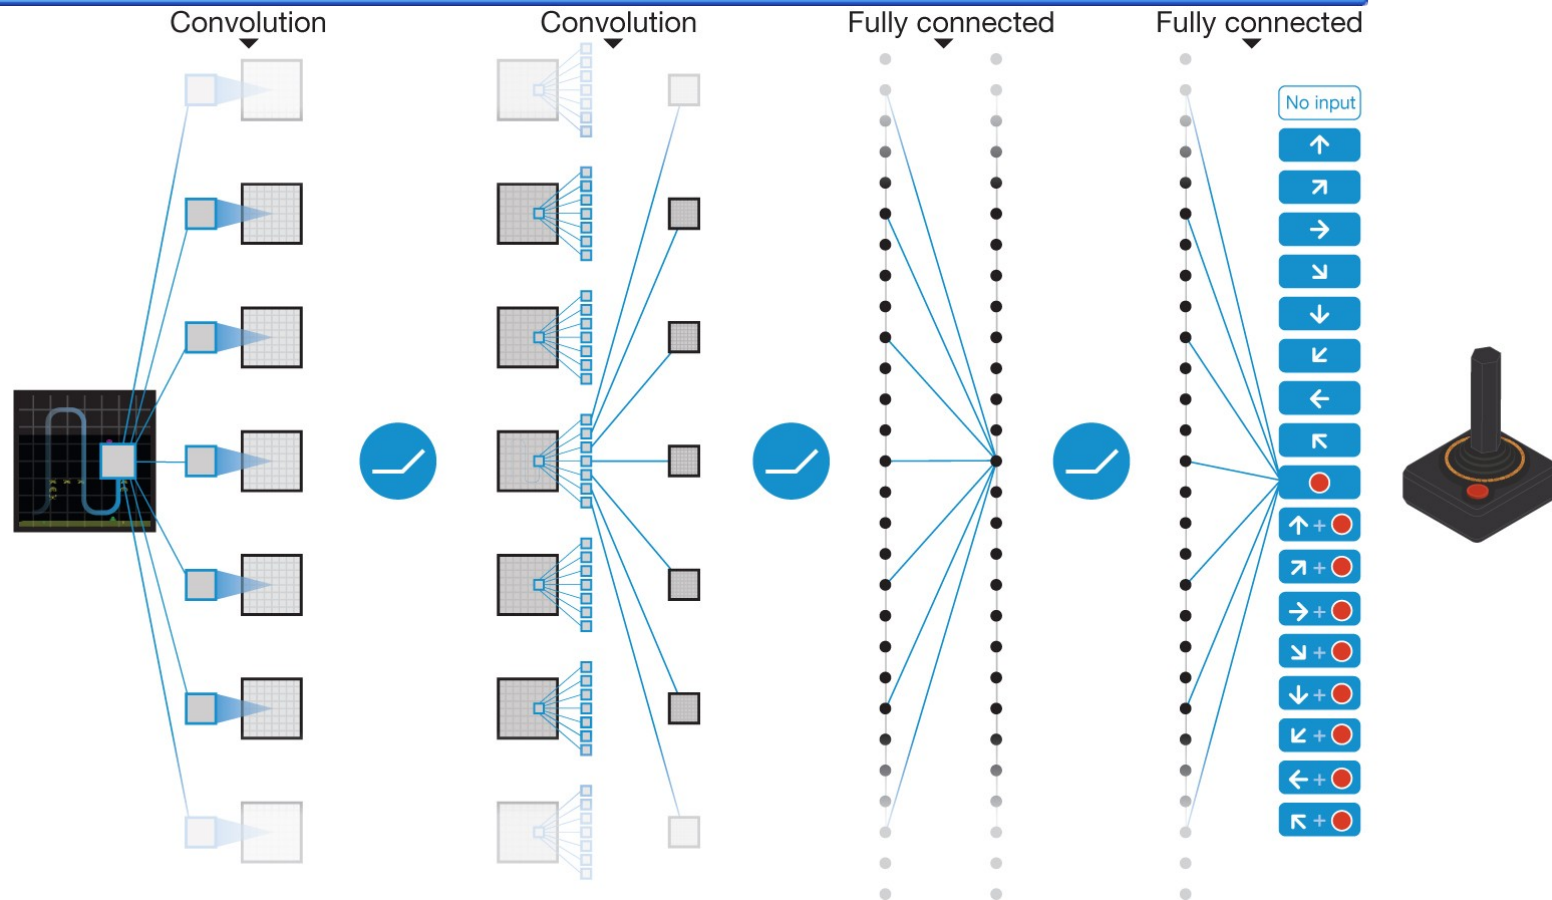
$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

$\gamma$  is the discount factor controlling agent's horizon,  $\theta_i$  are the parameters of the Q-network at iteration  $i$  also the network parameters used to compute the target at iteration  $i$ .

Target network parameters are only updated with the Q-network parameters ( $\theta_i$ ) every  $C$  steps and fixed between each update



# Schematic illustration (from Nature slides)



V Mnih et al. *Nature* **518**, 529-533 (2015) doi:10.1038/nature14236

# References

1. Human-level control through deep reinforcement learning by V Mnih, K Kavukcuoglu, D Silver, A Rusu, J Veness, M G Bellemare, A Graves, M Riedmiller, A Fidjeland, G Ostrovski, S Petersen, C Beattie, A Sadik, I Antonoglou, H King, D Kumaran, D Wierstra, S Legg, D Hassabis, Nature journal, Volume 545, pp. 529-533, February 2015, DOI:10.1038/nature14236

# References

2. The arcade learning environment: An evaluation platform for general agents by Marc G. Bellemare, Yavar Naddaf, Joel Veness, Michael Bowling, Journal of Artificial Intelligence Research 47, pages 253-279, DOI: 10.1613/jair.3912  
arXiv:1207.4708